

# Practical Concurrent Traversals in Search Trees

Tyler Townsend & Robert Bland

Authors:

Dana Draschler-Cohen,  
Martin Vechev,  
Eran Yahav



UCF

# Overview


- PaVT AVL & BST Data Structures
- Design of Data Structures
- GCC STM Data Structure
- Experimental Results
- Conclusion

# PaVT AVL & BST Data Structure

- Purpose of the paper is to propose a necessary and sufficient condition to validate traversals in Search Trees, PaVT.
- The authors then showed how to create a lock-free membership test that can be applied to any search tree using PaVT.
- The authors demonstrated PaVT in Binary Search Trees and AVL Trees.


# PaVT Motivation

- T1: Calls  
Contain(7)




# PaVT Motivation

- T1: Calls Contain(7)
- T1: Gets to node 8




# PaVT Motivation

- T1: Calls Contain(7)
- T1: Gets to node 8
- T2: Calls Remove(6)



# PaVT Motivation

- T1: Calls Contain(7)
- T1: Gets to node 8
- T2: Calls Remove(6)
- T2: Removes 6 and replaces with 7
- T1: Returns false



# Contributions over State-of-the-art

- PaVT (Path Validation in Search Trees)
  - Provides a condition to which a target node is or not in the tree.
  - Applies to any Search Tree
    - BST, Ternary Trees, 2-D Trees, Tries, ect..
- Observe a limited number of nodes that are the succinct path snapshot (SPS) of the path leading to the key.
- Concurrent Updates may modify the path and the traversal may still complete successfully if it has found the correct snapshot.



# PaVT with SPS in BST's

- Predecessor, Successor Pairs
- Similar to authors previous work of BST's with Logical Ordering
- When validating the path, check that the key may be reach from this path.
- Depends on whether the the last node traversed was left or right of parent.

```
if (parentIsLarger && key <=node.Predecessor or  
    !parentIsLarger && key >=node.Successor) {  
    restart;  
}
```

# Design of Data Structure

- Authors provide 4 Algorithms in Pseudocode
  - TraverseNLock(Traverse)
  - Insert
  - Remove
  - Contains (Implied from Traverse)
  - UpdateSnaps
- Up to Us
  - Tree Structure
  - AVL Operations
  - Snap Shots (mentioned in previous work)

# Node and Tree

```
class Node {  
    int data;  
    Lock;  
    Node left;  
    Node right;  
    Node parent;  
    atomic<Node> leftSnap;  
    atomic<Node> rightSnap;  
    int height;  
    bool mark;  
}
```

```
class Tree {  
  
    Node root;  
    atomic<Node> minSentinel;  
    atomic<Node> maxSentinel;  
  
    traverse(Node, data);  
    insert(data)  
    remove(data)  
    contains(data)  
    rebalance(data)  
}
```

# Locking Protocol

Each node has a lock and  $n_1$ 's lock is acquired before  $n_2$ 's lock if

1.  $n_2$  is reachable from  $n_1$  or
2.  $n_1$  and  $n_2$  are reachable via  $f_i$  and  $f_j$ , respectively, from their lowest common ancestor and  $i < j$

Lock from top to bottom, left to right ordering.

Locks (parents) acquired out of order are done optimistically and call for restart if failed

- Protocol is Livelock and Deadlock free since locks are acquired in same order by all threads
- Lowest node in tree is guaranteed to succeed in acquiring locks

# Snapshots

- Snapshots are updated by Traversing from mutated Node up the path
- Can be modified directly in the case of BST's
- Modified and updated within insert/remove directly
- Correctness
  - Must acquire locks before mutating path
- Linearizability
  - When snapshots are read.
  - Linearized just before linearization point of insertion/removal
  - This represents a consistent path and since the snapshot has not been updated yet this implies the other thread possibly modifying it has not completed its operation.

# Traverse

- Used by Insert and Remove
- Contains is exactly the same as traverse except we do not use locks.
- Correctness follows from PaVT and Linearizes when node is locked and returns

```
Traverse(node, data) :  
    n = node  
    f = nextField(node, data)  
    while(n.f != null) {  
        n = n.f  
        if node is found  
            if marked then restart  
            lock n and return n  
    }  
    check if data is in S(n,f)  
    if true restart  
    else lock n and return n
```

# Rebalance(node)

- Continually traverse up to root from node
- Acquire locks of parent, node, child, grandchild
- Check balance factors of node, child
- Perform single/double rotations when necessary
- If current height of node is unchanged and balance factor is satisfied then we are done
  - Based off relaxed AVL approach
- Linearizes after last rotation per traversal
- Must be very careful with new locking order!

# Insert(7)

Logical Ordering: 1, 4, 6, 12

Insert(data) :

**node = traverse(data)**

check if insert to null pointer

newNode = new Node(data)


updateParent(node, newNode)

setField(node, newNode)

Copy snaps of node into newNode

UpdateSnaps of node and oldSnap

Unlock(node)



Legend

Tree	_____
Snapshot	↔



# Insert(7)

Logical Ordering: 1, 4, 6, 12

Insert(data) :

node = traverse(data)

check if insert to null pointer

newNode = new Node(data)


updateParent(node, newNode)

setField(node, newNode)

Copy snaps of node into newNode

UpdateSnaps of node and oldSnap

Unlock(node)



Legend

Tree



Snapshot



# Insert(7)

Logical Ordering: 1, 4, 6, 12

Insert(data) :

node = traverse(data)

check if insert to null pointer

**newNode = new Node(data)**


**updateParent(node, newNode)**

**setField(node, newNode)**

**Copy snaps of node into newNode**

UpdateSnaps of node and oldSnap

Unlock(node)



Legend

Tree

Snapshot



# Insert(7)

Logical Ordering: 1, 4, 6, 7, 12

Insert(data) :

node = traverse(data)

check if insert to null pointer

newNode = new Node(data)



updateParent(node, newNode)

setField(node, newNode)

Copy snaps of node into newNode

UpdateSnaps of node and oldSnap

Unlock(node)



# Insert(Data)

- AVL implementation attempts to rebalance after unlocking if it is needed
- Insert linearizes when the node is added to the parent

# Remove(Data)

- "Five" possible Scenarios for a node being removed
  - a. Node has no children
  - b. Node has 1 leaf
  - c. Node has 2 children (let  $r$  = node right child)
    - $r$  has no left child
    - successor's parent is  $r$
    - successor's parent is not  $r$
- Biggest Challenge is preserving locking/unlocking order!

# Remove(7)

Logical Ordering: 1, 4, 6, 7, 12

Remove(data) :

**node = traverse(data)**

parent.try\_lock()


mark(node)

setField(Parent, null)

updateParent(parent, rightSnapshot)

Change right snap of parent to child's  
successor

UnlockAll()



Legend

Tree



Snapshot



# Remove(7)



Logical Ordering: 1, 4, 6, 7, 12

Remove(data) :  
node = traverse(data)  
**parent.try\_lock()**  
mark(node)  
setField(Parent, null)

updateParent(parent, rightSnapshot)

Change right snap of parent to child's  
successor

UnlockAll()



# Remove(7)



Logical Ordering: 1, 4, 6, 7, 12

Remove(data) :  
node = traverse(data)  
**parent.try\_lock()**  
mark(node)  
setField(Parent, null)

updateParent(parent, rightSnapshot)

Change right snap of parent to child's  
successor

UnlockAll()





# Remove(7)



Logical Ordering: 1, 4, 6, 7, 12

Remove(data) :  
node = traverse(data)  
parent.try\_lock()  
**mark(node)**  
setField(Parent, null)

updateParent(parent, rightSnapshot)

Change right snap of parent to child's  
successor

UnlockAll()



# Remove(7)

Logical Ordering: 1, 4, 6, 7, 12

## Remove(data) :

```
node = traverse(data)
```

```
parent.try_lock()
```



mark(node)

```
setField(Parent, null)
```

```
updateParent(parent, rightSnapshot)
```

## Change right snap of parent to childs successor

## UnlockAll()



# Remove(7)



Logical Ordering: 1, 4, 6, 12

```
Remove(data) :  
  node = traverse(data)  
  parent.try_lock()  
  mark(node)  
  setField(Parent, null)
```

```
updateParent(parent, rightSnapshot)
```

Change right snap of parent to child's  
successor

UnlockAll()



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

**node = traverse(data)**

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()

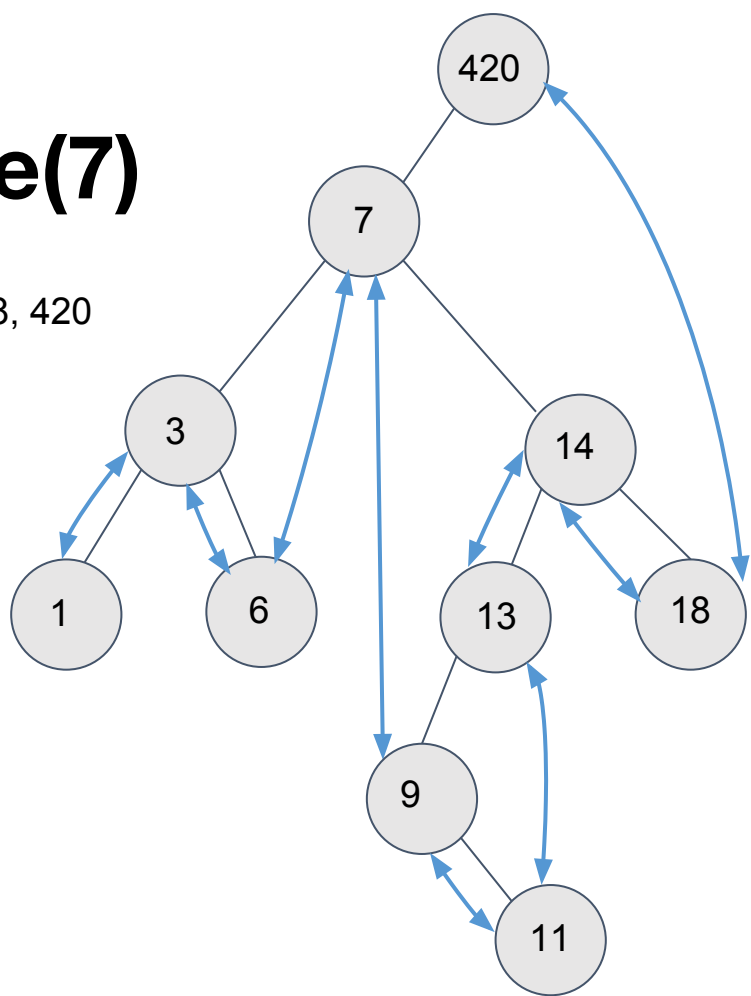
succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

**node = traverse(data)**

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()

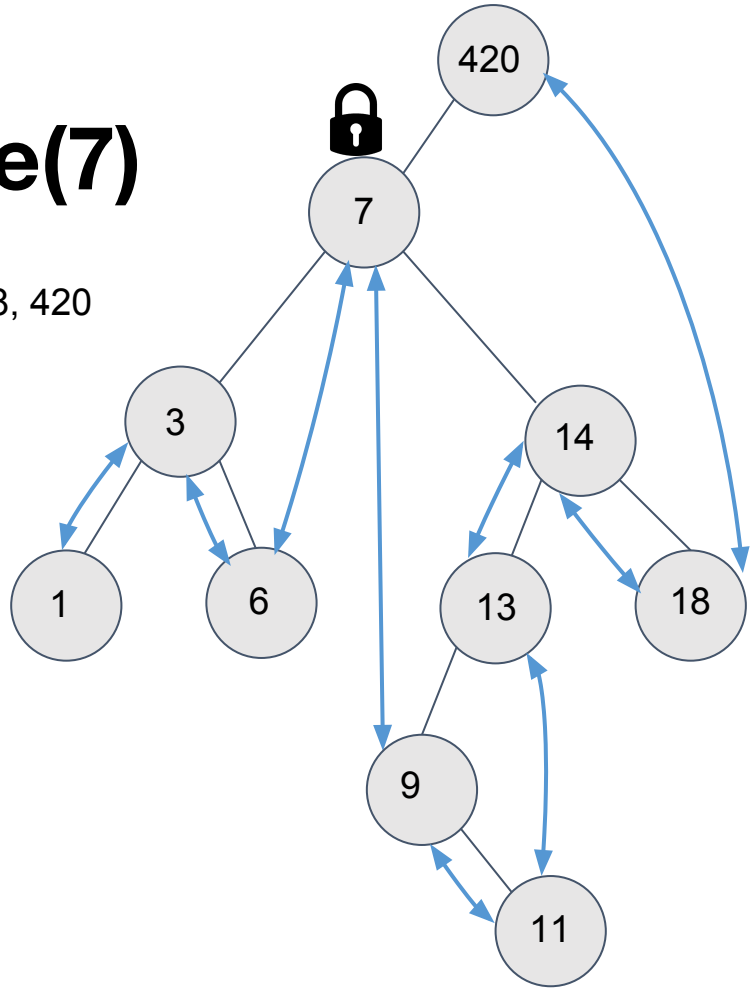
succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()

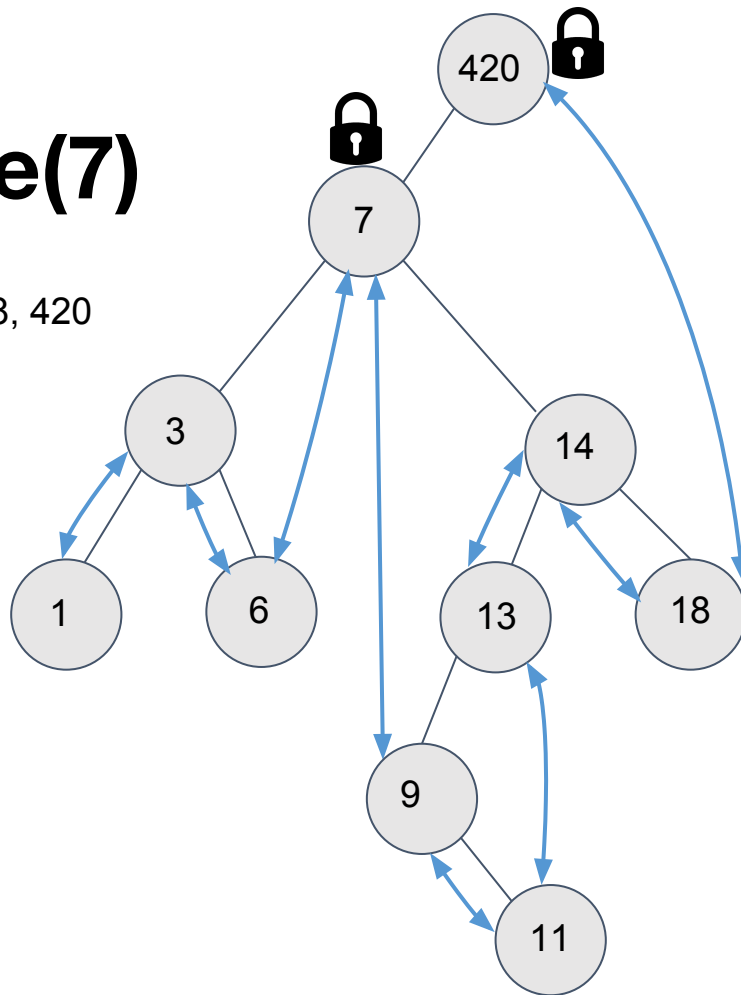
succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

**rightChild.lock()**

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot





# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



UCF

# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

**succ.lock()**


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


**succChild.lock()**

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree



Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree

Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 7, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

**bstRemove()**

updateSnaps(pred, succ)

unlockAll()



Legend

Tree

Snapshot



# Remove(7)

Logical Ordering :

1, 3, 6, 9, 11, 13, 14, 18, 420

Remove(data):

node = traverse(data)

parent.try\_lock()

leftChild.lock()

rightChild.lock()

pred.lock()

succParent.lock()

succ.lock()


succChild.lock()

mark(node)

bstRemove()

updateSnaps(pred, succ)

unlockAll()



Legend

Tree

Snapshot



# Remove(data)

- Deletes the node from the tree and replaces it if the node being deleted is not a root
- Uses locks by calling traverse
- Uses additional locks based on whether or not the node being removed has children
- Attempts to lock against locking order when locking the parent so only a trylock is used
  - If this fails then we restart
- Attempts to rebalance once removal is finished if it is an AVL tree
- Remove linearizes when it successfully marks the removed node



# Reimplementation

- We faced two main challenges with our implementation
  - Creating and maintaining snapshots
    - Snapshots were a bit obfuscated in the paper
    - The authors were not precise on what they represented in their example
  - Handling the worst case scenario of Remove
    - The worst case scenario involves 8 locks
    - This made handling and debugging them a terror when other scenarios only had to lock a few nodes
    - No mention of locking the predecessor of the node being removed

# GCC STM Data Structure

- GCC Transactional memory is used to create blocks of code that will be executed atomically
- Important Method/Properties of GCC
  - `__transaction_atomic{}`
    - All code in this statement block will be executed atomically
    - By default it executes in an all or nothing fashion and will repeat until it successfully commits
    - Will not execute transaction unsafe code.
  - `transaction_safe` & `transaction_unsafe`
    - A property that can be applied directly to functions
    - By default the GCC can determine `transaction_safe` functions
    - However, this is only true if the code is defined in the source file

# STM Implementation

The steps we took to create the GCC implementation are:

- 1) Remove atomic wrappers from objects as they are not transactionally safe
- 2) Mark methods outside the source file as transactionally safe
- 3) Atomic Transactions:

```
insert(data) {  
    __transaction_atomic {  
        node = traverse(root, data)  
        ...  
        if (avl) rebalance(node)  
    }  
}
```

```
remove(data) {  
    __transaction_atomic {  
        node = traverse(root, data)  
        ...  
        if (avl) rebalance(node)  
    }  
}
```

# STM Implementation

Contains only requires atomic reads for the snaps shots so it still provides concurrency!


# Improved GCC STM Implementation

- Decouple two functions from the atomic block
- Traverse
  - Snapshots are now checked inside of traverse
  - Atomic blocks contain  $O(1)$  write operations, no longer traversing
- Rebalancing
  - Decoupled rebalancing from all functions where they are called
  - Rebalancing has its own atomic transaction
  - Thread will allow others to operate while it attempts to rebalance the tree.
  - Traversal is done outside of atomic updates

# Experimental Results


- Experimental Details
  - Intel(R) Core(TM) i5-4460 CPU @ 3.2GHz
  - Number of Cores/Threads 4/4
  - 24.0 GB RAM
  - Ubuntu 18.04
- BenchMarks (Insert /Remove /Contains):
  - 9% /1% /90%
  - 20% /10% /70%
  - 50% /50% /0%
- Tree initially filled to expected capacity (90%, 2/3, 50%)

# Data Structure Results



- AVL generally outperforms BST except where there are high number of inserts/removes
  - This scenario would force more rebalancing and AVL would have less leafs as a result
- Scaled linearly until we hit a wall at four threads which we suspect is due to hardware limitations

# STM Results




- AVL outperformed BST on most occasions except when it comes to high inserts/removes
  - This is due to the amount of rebalancing that gets called that would interfere with all transactions
- Performance sharply drops in all cases due to hardware




# STM 2.0: Results


STM 2.0: 9%/1%/90%



STM 2.0: 20%/10%/70%



STM 2.0: 50%/50%/0%



- In the first graph the output for BSTs increased by 50 percent with one and two thread and by about 10 percent with four threads
- In the second graph the output for BSTs increased by 50 percent up to 4 threads
- AVL has a slight increase overall with the notable being the last graph having a 50% increase with four threads

# Conclusion

- Our implementation scales well until we hit our hardware limitation
- Transactional implementation performs worse than our original data structure in most scenarios except when there are a high number of inserts & removes
- PaVT is show to provide lock free membership tests.
- Snapshots are easy to maintain and implement for the search trees and take up very little space